
Reinforcement Learning in an Asymmetric, Perfect-Information Historical Wargame: The Yom Kippur War

Renato Vicente

Department of Applied Mathematics
Institute of Mathematics, Statistics and Computer Science
University of São Paulo
Rua do Matão, 1010, 05508-090 São Paulo, SP, Brazil

Abstract

We present a complete reinforcement-learning (RL) case study on a historical hex-and-counter wargame — *The Yom Kippur War* — digitally reconstructed from its original components (board, counters, and combat tables). The game has perfect information with randomness confined to dice-based combat resolution (2d6), short episodes (6 rounds, ~ 90 micro-decisions per side), and strong **asymmetry** between an attacker (Israel), who must chain a long-range amphibious operation, and a defender (Egypt), which is reactive and locally advantaged. We build a vectorized environment on the game’s own rules engine, with a masked action space decomposed into micro-decisions, and train PPO policies with three architectures: an MLP, a CNN over the hexagonal grid and a GNN over the board’s exact topology. An architecture-agnostic interpreter serves them in the browser with verified numerical parity. The study then adds PUCT search with a learned value, an AlphaZero-style loop and temporal abstraction with an LLM strategist. Results on architecture–role fit and on the collapse of naive self-play were measured against the heuristic opponent of the time and are reported as such. **A later audit of the evaluation reverses the study’s central negative claim.** The heuristic opponent changed during the study and most comparisons rested on 12 games, so the apparent ceiling near 33% for the attacker was an artifact. Retrained against the current defender, the attacker wins 87% instead of 14%; adding the heuristic’s move to the candidates of PUCT raises search from 23.5% to 55%; and tuning the heuristic’s weights by the cross-entropy method nearly triples the heuristic attacker. Every agent trained against a single opponent, however, is exploited by some counter-strategy: the 87% attacker loses all 1000 validation games to a reweighted heuristic defender. **We therefore summarize the game by the Nash equilibrium of a population** of eight attackers and six defenders after one PSRO iteration, in which the attacker wins 17%. We release the game, environment, trainer, search, evaluation arena and models as a reproducible artifact.

1 Introduction

Board games have served for decades as testbeds for sequential decision-making. Operational hexagonal wargames (*hex-and-counter*), however, are comparatively underexplored by RL despite attractive properties: discrete, well-defined rules; a short horizon; and — in the game studied here — **structural asymmetry** between the two sides, which exposes multi-agent RL difficulties that are rarely visible in symmetric games such as Go or Chess.

We study *The Yom Kippur War* (Wargame II), a simulation of the Israeli crossing of the Suez Canal in October 1973. The Israeli general (attacker) has 6 rounds to destroy surface-to-air (SAM) batteries, establish a bridgehead on the west bank, and control city entrances; the Egyptian general (defender) seeks to prevent this. The historical operation was notably audacious precisely because the position favored the defender — an asymmetry that is mirrored in the RL dynamics.

Contributions.

1. An **RL environment faithful to the production rules engine** (the same code that runs the playable game), eliminating train/deployment divergence.
2. A **micro-decision action space with a legality mask** derived from the engine itself, making a combinatorial turn (moving up to 19 units) tractable.
3. An **architecture-agnostic layer interpreter** that lets us train MLP/CNN/GNN in PyTorch and serve any of them in pure JavaScript in the browser, with tested numerical parity.
4. An **AlphaZero-style PUCT search with a learned value** (opponent and chance folded into the transition, leaves scored by the PPO critic) and a complete self-play retraining loop, both runnable in the browser/Node from the same exported networks.
5. Empirical findings on **architecture $\leftrightarrow$ role fit, self-play collapse in asymmetric games** and its **partial mitigation**.
6. An **audit of the evaluation** (Section 7): a parallel arena with Wilson intervals, the finding that the opponent changed mid-study, retraining and cross-entropy tuning of both sides, a candidate-set fix for search, and a **population Nash** after one PSRO iteration as the summary of the game.

2 Related Work

Tree-search self-play methods (AlphaZero and successors) reached super-human performance on symmetric perfect-information board games. PPO is an established, robust baseline for policy-based control. In multi-agent RL, naive self-play is known to be unstable (non-transitive cycles, catastrophic forgetting), motivating opponent leagues and *fictitious play*. Potential-based reward shaping provides dense signal without changing the optimal policy. KL penalties against a reference policy are used to limit drift and preserve knowledge. Our work proposes no new algorithm; it combines known techniques in an underexplored asymmetric domain and documents the resulting phenomena, with an emphasis on reproducibility and deployment.

3 The Game as an RL Environment

3.1 Domain

The board is a hexagonal grid of **538 valid cells** (columns 1–25, rows 1–22; flat-top topology with even columns shifted), with **1521 adjacency edges**. Seven terrain types modulate movement cost and defensive bonus; roads, the Suez Canal, and a freshwater canal are edge features. Each side controls battalions with a Combat Index (CI) and a Mobility Index (MI); units take casualties (flip to a reduced-value back) and are eliminated on the second casualty.

A round has **8 phases** (movement, attack declaration, defensive support, and resolution — first for Israel, then for Egypt). Combat between adjacent units is **mandatory**; the outcome follows the force ratio (with terrain) mapped to a column (−2 to +2) of an Effects Table read with 2d6. The only source of randomness is this roll. Victory conditions are asymmetric and graded (decisive/partial/marginal) per side.

3.2 MDP formulation

- **Episode:** one game (6 rounds).
- **Decision (step):** one micro-decision — a unit’s destination, a reinforcement’s entry cell, or an artillery support target.
- **Action:** an integer in $\{0, \dots, 537\}$ (cell index) $\cup$ {PASS}, totaling **539** actions. The cell’s semantics depend on the current decision.




Figure 2: System pipeline. The same rules engine backs both the playable game and the vectorized training environments; trained policies are exported as an architecture-agnostic layer spec and served in-browser by a pure-JS interpreter.

Doctrine via target lists. High-level behavior emerges from per-unit target sets rather than a planner. For Israel: the engineer seeks the best bridge-eligible cell (east bank, between the lakes, far from enemies); before the bridge is up the maneuver units escort/rescue it; afterwards their targets become the uncontrolled victory objectives (Ismailia/Fahid entrances, live SAM cells). For Egypt: units already on a defensive objective garrison it; if the Israeli bridge is up, the engineer becomes the army’s priority target (destroying it removes the bridge and cuts reinforcements); otherwise units spread between objectives and containment.

Position scoring. In the movement phase, the heuristic enumerates a unit’s reachable cells (from the engine’s terrain-/ZOC-aware Dijkstra) and picks the arg max of a hand-crafted scalar score. For a candidate cell c and unit u ,

$$\text{score}(u, c) = -10 d(c, \mathcal{T}_u) + 80 \mathbf{1}[c \in \mathcal{T}_u] + \rho(\kappa(u, c)) + 2 b(c) - \xi(u, c),$$

where $d(c, \mathcal{T}_u)$ is the hex distance to the nearest target in the unit’s target set $\mathcal{T}_u$; $b(c)$ is the terrain defensive bonus; and ξ penalizes *exposure* (much more enemy than friendly Combat Index within radius 3). The dominant term $\rho(\kappa)$ reflects that combat is *mandatory*: ending adjacent to an enemy commits the unit to attack next phase, so the heuristic estimates the resulting Combat-Results-Table column κ and scores

$$\rho(\kappa) = \begin{cases} +70 & \kappa \geq +2 \quad (\text{triple-odds attack, no elimination risk}) \\ +35 & \kappa = +1 \\ -25 & \kappa = \text{“}\pm\text{”} \quad (\text{eliminates the attacker on high rolls}) \\ -100 & \kappa < 0, \end{cases}$$

with an extra -80 for artillery/engineer (support units should not engage). The **column estimate** $\kappa(u, c)$ simulates the combined combat the move would trigger: it sums the Combat Index of the moving unit with that of all friendly units already adjacent to the affected enemies (which the rules fuse into a single combat), adds each defender’s terrain bonus, and maps the resulting force ratio to a column in $\{-2, \dots, +2\}$. Artillery barrage/cover is allocated only when it *improves* that column (each battery fires once per round); resolution-time choices sacrifice the weakest/wounded unit and retreat so as to maximize distance from the enemy.

Two scoring modes. The fixed weights $\rho(\kappa)$ were in use when the curricula of Section 6.1 were trained. A later variant replaces ρ by the expected utility of the combat, $65 \mathbb{E}_{2d6}[U(\kappa)]$, where U assigns a utility to each outcome of the Combat Results Table and the expectation runs over the 2d6 distribution. This *expectimax* mode became the default for both sides and is the opponent in every evaluation from Section 6.4 onward. As Section 7.1 shows, the change is far from neutral: it strengthens the defender considerably and weakens the attacker.

Policy definitions. A neural policy defines $\pi_\theta(a \mid s) = \text{softmax}(\text{mask}(z_\theta(s)))$, where z_θ are the per-action logits and $\text{mask}(\cdot)$ sets illegal actions to $-\infty$; we sample $a \sim \pi_\theta$ during training (for exploration) and take $\arg \max_a \pi_\theta$ at evaluation (greedy). The heuristic policy is the deterministic $\arg \max$ of $\text{score}(u, \cdot)$ above. Both expose the identical interface $\text{playPhase}(s)$ over the engine, so any side of any game can be driven by either a network, the heuristic, or a human, interchangeably — which is what lets the same trained policy serve as a curriculum opponent, a league member, and an in-browser AI.




Figure 3: The three policy/value trunks. The MLP treats the observation as a flat vector; the CNN convolves over the offset hex grid (each cell scattered to a 22×25 lattice); the GNN performs message passing over the exact 1521-edge hex adjacency. CNN and GNN share weights across space and are two to three orders of magnitude smaller than the MLP. “GNN msg” concatenates each node’s features with the mean of its neighbors.

4.2 Architectures

Three policies with the common interface $(\text{obs}, \text{mask}) \rightarrow (\text{logits}[539], \text{value})$. The trunk differs; both heads are shared in spirit: a *policy head* emits one logit per cell plus a PASS logit, and a *value head* maps a global pooling (concatenated with scalars) to a scalar. The mask adds $-\infty$ to illegal logits before the softmax. Figure 3 details the three trunks.

4.3 Training (PPO) and deployment

We train with PPO using advantage estimated against a value baseline, clipping, and an entropy bonus. Policies are exported as a **list of executable layers** (`linear`, `conv`, `gnn`, `heads`), interpreted both by training and by a pure-JavaScript inference engine — enabling play against any architecture in the browser with no dependencies. A parity test confirms that the interpreter’s 539 logits match PyTorch (max error $< 10^{-8}$).

4.4 Self-play and stabilization

In self-play both sides are policies in training, with a **league** of frozen checkpoints (*fictitious play*) and per-trajectory credit ($\gamma = 1$). To curb the attacker’s degradation (Section 6.2), we combine:

- **Potential-based shaping:** $\Phi(s)$ grows with attacker progress (bridge, SAMs, objectives, bridge-head). With $\gamma = 1$ the return telescopes to $R_{\text{terminal}} + \alpha (\Phi_{\text{final}} - \Phi_t)$, preserving the optimum (Ng et al., 1999).
- **Asymmetric self-play:** more epochs, learning rate, and fraction of environments for the weak side.
- **KL penalty against an anchor:** $\beta \text{KL}(\pi_{\text{cur}} \parallel \pi_{\text{anchor}})$, the anchor being the attacker’s curriculum policy.

4.5 Search at inference: PUCT with a learned value

At play time we optionally replace greedy policy evaluation with an AlphaZero-style Monte-Carlo tree search, exploiting two structural simplifications of the domain. First, because combat is mandatory and the opponent’s phases and the 2d6 Effects Table are part of the environment, we **fold the opponent and chance into the transition**: between two of our movement decisions the heuristic plays the opponent’s round and the engine resolves combat (using the expectimax column choice of Section 4.1). The tree then contains only *our* decision nodes, the value is always from our perspective, and no minimax sign-flipping is needed — a single-agent search against a fixed world model. Second, we **never roll out to the end**: leaves are scored by the network’s value head V_θ , the PPO critic trained to predict the graded terminal outcome, which supplies the long horizon a 1-ply policy lacks. Selection uses PUCT,

$$a^* = \arg \max_a Q(a) + c_{\text{puct}} P(a) \frac{\sqrt{\sum_b N(b)}}{1 + N(a)},$$

with prior P the masked, top- k policy head and Q the mean backed-up leaf value; the chosen move is the most-visited root action. Only the movement phase is branched; combat phases use the base policy. Defaults: $c_{\text{puct}}=1.5$, $k=6$, 48 simulations, depth 4 of our rounds. The value head, already exported with every network, is surfaced to the JS interpreter; the parity test of Section 4.3 is extended to confirm both logits *and* value match PyTorch to $< 10^{-8}$.

AlphaZero refinement. We further close the learning loop: PUCT self-play generates targets (π, z) , where π is the root visit distribution (with root Dirichlet noise and an early-move temperature schedule for exploration) and z the graded outcome, on which we retrain the policy (cross-entropy to π) and value (MSE to z), warm-started from the PPO models, and iterate. We also study a *value-only* variant that freezes the policy and refits only V_θ , motivated by the result below.

5 Experimental Setup

Phase-A curricula train each side against a fixed heuristic AI for 150–200 updates (32 parallel environments). Evaluation is greedy, over 24–100 games with held-out seeds, reporting win rate by type. For self-play, we seed the policies with the best curricula. Main hyperparameters: $\text{lr } 3 \cdot 10^{-4}$ (curriculum), clip 0.2, $\lambda_{\text{GAE}} 0.95$, entropy 0.01–0.03. Fixed seeds make games reproducible.

Two evaluation protocols. Section 6.1 through Section 6.6 report the protocol used when those experiments were run: greedy play over 12 to 100 games, against the heuristic in whichever scoring mode was current. Section 7 re-evaluates with a parallel arena that plays complete games across workers, reports Wilson 95% intervals [7], and uses at least 300 games per condition and 1000 to validate optimized agents, always on seeds that were not seen during training or tuning. Where the two disagree, the second should be preferred.

6 Results

6.1 Architecture $\leftrightarrow$ role fit (vs. heuristic)

Side	MLP (2.86 M)	CNN (18 k)	GNN (6.7 k)
Egypt (defender; reactive, local)	100%	—	92%
Israel (attacker; long-range plan)	97%	0%	0%

The Egyptian side, defensive and locally decided, is comfortably solved by small, spatially structured networks — the GNN reaches 92% with $400\times$ fewer parameters than the MLP (~ 35 KB vs. ~ 15 MB model). The Israeli side, which must chain bridge $\rightarrow$ crossing $\rightarrow$ capture of scattered objectives, **is mastered only by the high-capacity MLP**; CNN and GNN converge prematurely (entropy $\rightarrow 0$ around update 100) to a local optimum of “losing by a smaller margin” (reward ~ -0.6), never winning. Section 6.5 traces this specifically to *capacity* (a wider spatial net does win), not to an inductive-bias mismatch as one might first assume. These rates were measured against the fixed-weight heuristic. Against the expectimax defender used in later sections, the same attacker MLP wins 14.0% (Section 7.1).

Table 1: Self-play stabilization. ISR-vs-heuristic “= 1 rate” is the fraction of greedy evaluations in which the attacker still wins 100% against the heuristic — a proxy for not having forgotten its curriculum competence. Head-to-head is the attacker’s win rate against the co-trained defender.

	Naive	Mitigated ($\beta=0.05$)	Aggressive ($\beta=0.10$)
ISR vs. heuristic (mean)	$\rightarrow 0$ (stuck)	oscillates, recovers	0.83
ISR vs. heuristic (= 1 rate)	0	—	28/50
ISR×EGY head-to-head (peak)	~ 0	0.58	0.54
ISR×EGY head-to-head (mean)	~ 0	—	0.24
EGY vs. heuristic	strong	strong	0.71–0.96
$KL(\pi \parallel \text{anchor})$	—	≤ 0.39	0.10–0.22

6.2 Naive self-play collapse

Starting from the curriculum policies (ISR 97%, EGY 92–100% vs. heuristic), unmitigated self-play **monotonically degrades the attacker**: its rate vs. the heuristic falls from 97% to 0% and the head-to-head ISR×EGY sits at ~ 0 . The structurally favored defender wins almost always; with no positive reward signal, the attacker’s gradient drives it away from any prior competence (catastrophic forgetting).

6.3 Stabilization

We combine potential-based shaping, asymmetric self-play, and a KL anchor against the attacker’s curriculum policy. Table 1 compares naive self-play with two mitigation settings: a moderate one ($\alpha=0.1$, $\beta=0.05$, 200 updates) and an aggressive one ($\alpha=0.15$, $\beta=0.10$, 8/2 epochs and 70% of environments for the attacker, 250 updates).

The mitigations **prevent the permanent collapse** of the naive case (where the attacker is stuck at 0%): the KL penalty re-anchors the policy whenever it degrades, so the attacker spends most of training at 100% vs. the heuristic (mean 0.83; perfect in 28 of 50 evaluations) and gains real head-to-head competitiveness (recurring peaks of 0.46–0.54, against ~ 0 for naive self-play). The aggressive setting is the more stable of the two. **Full balance is not reached**, however: the head-to-head mean stays below 0.5 because the defender is structurally favored. The techniques solve the stated problem (avoiding collapse) but do not make the game symmetric — a limit of the domain, not of the method.

6.4 Search strengthens the attacker; the value dominates the prior

We now add inference-time search (Section 4.5). Table 2 compares, apples-to-apples (same 12 held-out seeds, same full game loop), the heuristic, the greedy neural policy, and PUCT. This protocol is stricter than the curricula of Section 6.1: absolute rates are lower because the seed set strongly favors the defender (heuristic-vs-heuristic gives ISR 8% / EGY 92%); the *relative* ordering is the point. **PUCT quadruples the attacker** (8% $\rightarrow$ 33%) and doubles the greedy policy (17% $\rightarrow$ 33%): the value head delivers the strategic horizon the 1-ply policy lacks (it now “sees” the payoff of bridge $\rightarrow$ crossing $\rightarrow$ capture). The defender, already near its heuristic ceiling (92%), is not improved (PUCT and greedy both 67%): search cannot exceed an evaluator that is itself weaker than the heuristic on defense.

Note added in revision. The explanation of the low absolute rates given above is incorrect. They reflect a change of opponent, the expectimax defender of Section 7.1, rather than the seed set. On 300 games the heuristic baseline is 15.3%, not 8%, and with 12 games each rate in Table 2 carries a 95% interval of about ± 25 points.

A control isolates *what* the search relies on. Driving PUCT with a *stronger* policy prior but a worse-calibrated value — the 97% curriculum MLP, whose critic only ever saw heuristic play — *lowers* the attacker to 17%, below the 33% obtained from the self-play model’s weaker policy but better-calibrated critic. **In PUCT here the value dominates the prior**: search quality tracks the evaluator, not the actor. This single observation governs all of the AlphaZero results that follow.

Table 2: Inference-time search vs. greedy play and the heuristic (full game loop, 12 defender-favoring held-out seeds, win rate vs. the heuristic). PUCT uses 48 simulations.

Side	Heuristic	Greedy policy	PUCT (48 sims)
Israel (attacker)	8%	17%	33%
Egypt (defender)	92%	67%	67%

Table 3: AlphaZero refinement (win rate vs. heuristic, 10 held-out seeds), warm-started from the PPO models. The value-only variant freezes the policy and refits only V_θ .

Variant	Israel	Egypt
Baseline (PUCT, PPO value)	30%	70%
Naive AZ (policy + value)	0%	60%
Value-only, lopsided self-play	0%	80%
Value-only, balanced targets	10%	—

AlphaZero refinement: an instructive asymmetry. Table 3 reports three controlled variants of the loop (Section 4.5, 10 seeds). (i) Naive joint retraining **degrades** both sides: the sharp PPO prior collapses toward the near-uniform π produced by a low-simulation search. (ii) Freezing the policy and refitting *only* the value isolates the cause — with the prior intact, the **defender improves** (70 $\rightarrow$ 80%) because its self-play outcomes are informative, but the **attacker collapses to 0%**: in defender-dominated self-play almost every attacker record is a loss ($z < 0$), so its value head loses the ability to *discriminate* good positions from bad. (iii) Regenerating the attacker’s targets against the heuristic (balanced, $\sim 42\%$ wins) recovers only part of the gap (0 $\rightarrow$ 10%, still below the 30% PPO baseline): raw Monte-Carlo targets from 24 games are noisier than the PPO critic (trained with bootstrapping over thousands of games), and the narrow-margin attacker is sensitive to value variance. Surpassing the PPO value via AlphaZero would require large-scale data and bootstrapped value targets, infeasible at ~ 10 s per searched game on one machine. The pipeline is released ready to scale.

6.5 Revisiting the small-network attacker: capacity, not inductive bias

Section 6.1 left open *why* only the high-capacity MLP masters the attacker while the CNN/GNN converge to a “lose-by-less” optimum (reward ~ -0.6 , 0% wins). A natural reading is an inductive-bias mismatch (translation-equivariant weight sharing cannot pinpoint the single bridge cell or named objectives). A controlled ablation **refutes** that explanation. Curriculum training is sparse-reward (graded terminal only); we first add the potential-based shaping of Section 4.5 (Φ : bridge, SAMs, objectives, bridgehead) to the attacker’s curriculum, then vary network width (Table 4). Shaping alone does *not* help a small net: the CNN/GNN never make partial progress, so Φ stays ≈ 0 and the shaping term is inert (0%, identical to unshaped). But a $7\times$ **wider GNN** (47.7k vs. 6.7k parameters) under the *same* shaped settings *learns to win*, peaking at 33% — on par with the MLP under this protocol — while the narrow GNN with identical shaping stays at 0%. The lever is therefore **capacity, not inductive bias**. The wide net is, however, **unstable**: with the default entropy bonus it oscillates (peaks at 33%, then forgets to 0–4%, entropy pinned near its maximum). Lowering the entropy coefficient (0.02 $\rightarrow$ 0.005) and learning rate lets it commit (entropy $\rightarrow$ 1.5) and **sustain** ~ 15 –21%. The small-network attacker failure is thus one of *capacity and optimization stability*, not of representational bias.

6.6 Temporal abstraction and an LLM strategist: a hierarchy law

The attacker’s difficulty is long-horizon (~ 90 micro-decisions per game). We test *temporal abstraction*: a high-level controller chooses, once per round, among a few **macro-actions** (which objectives to commit force to), and a low-level executor realizes them. The macro-action reuses the heuristic’s target-set mechanism, so the existing movement engine is the executor for free; leaves are scored by the value head. This collapses the branching from ~ 90 cells to ~ 5 plans per round. We also study an **LLM strategist** (the hybrid “LLM proposes, simulator/value verifies”): a frozen language model is given a textual board encoding and returns candidate plans in JSON, which are *grounded* (illegal

Table 4: Attacker (Israel) curriculum vs. heuristic with potential-based shaping, varying the GNN width. Width — not shaping — unlocks winning play; lower entropy stabilizes it.

Setting (GNN, shaped curriculum)	Params	Win vs. heur.
Narrow, entropy 0.02	6.7 k	0% (all evals)
Wide, entropy 0.02	47.7 k	peak 33% , unstable
Wide, entropy 0.005, lower lr	47.7 k	~ 15–21% sustained

Table 5: Temporal abstraction / LLM strategist for the attacker (win rate vs. heuristic, 12 held-out seeds). The macro layer helps the weak heuristic executor but not the strong PUCT executor. Plan generation by GPT-4o-mini / GPT-5 (OpenAI) and GPT-4o for the hierarchy.

Attacker policy	Win vs. heuristic
Heuristic	8%
Greedy neural policy	17%
Macro-actions, fixed vocab (heur. executor)	17%
LLM plans, GPT-5 (heur. executor)	17%
LLM plans, GPT-4o-mini (heur. executor)	25%
PUCT executor, 24 sims (no macro)	25%
Hierarchy: GPT-4o plan + PUCT, 24 sims	17%
Plain PUCT, 48 sims	33%

target cells dropped) and selected by the value head. Finally, a **true two-level hierarchy** keeps the LLM strategist on top but replaces the heuristic executor with the PUCT searcher of Section 4.5, biased toward the chosen plan’s targets.

Table 5 reports attacker win rate (vs. heuristic, 12 held-out seeds). Macro-actions over the *heuristic* executor double it (8% → 17%) and the LLM’s adaptive plans add more (→ 25%), confirming that temporal abstraction and a stronger strategist both help *when the executor is weak*. But a stronger generator does not: GPT-5 (17%) matches GPT-4o-mini (25%) within one game, and lookahead depth saturates (macro depth 2=3) — the heuristic executor is the ceiling. Crucially, putting the same macro plan *on top of the strong executor* (PUCT) does **not** help: at matched simulations the plan bias moves PUCT from 25% to 17% (a one-game, within-noise change, but no gain), while plain PUCT is best (33% at 48 sims).

We summarize this as a **hierarchy law for this domain: the value of the strategic layer is inverse to the strength of the executor**. A coarse macro plan rescues a weak, myopic executor (the heuristic), but imposed on a value-aware tactical searcher it *distorts* an already-good prior, pulling units into tactically worse positions the unbiased search avoided. The LLM hybrid is a clean instance of grounded generate-and-verify (no illegal moves reach the board), and its pipeline doubles as the optimization target for text-space skill optimizers; but in this game it does not surpass search. The strongest attacker remains plain PUCT with a learned value.

Note added in revision. Every cell of Table 5 comes from 12 games, so the ordering it shows is not statistically established; we keep the hierarchy law as a hypothesis. The table also predates the finding of Section 7.3: a policy retrained against the same defender wins about 87%, well above every entry here.

7 Revisiting the Evaluation

A later audit of the evaluation protocol changed three of the conclusions above. The rebuilt evaluator plays complete games in parallel, reports Wilson intervals and uses at least 300 games per condition. This section reports what changed and why, and then extends the study to both sides.

7.1 A moving opponent and small samples

The heuristic opponent did not stay fixed during the study. We introduced the expectimax scoring of Section 4.1 after training the curricula of Section 6.1, and it became the default for both sides.

Table 6: The heuristic defender changed during the study. Israeli win rate with Wilson 95% intervals. In the first column both sides are heuristic; in the second the attacker is the Phase-A MLP of Section 6.1 and only the defender changes.

Heuristic scoring	Heuristic vs. heuristic ($n = 300$)	Phase-A MLP vs. heuristic ($n = 150$)
Fixed weights ρ (curriculum era)	29.7% [24.8–35.1]	86.7% [80.3–91.2]
Expectimax (all later evaluations)	15.3% [11.7–19.8]	14.0% [9.3–20.5]

Table 7: Attacker retrained against the expectimax defender (MLP, warm-started from Phase A). Israeli win rate on 400 held-out games (150 for the first row).

Attacker checkpoint	Win vs. current defender
Phase-A MLP, not retrained	14.0% [9.3–20.5]
No shaping, update 30	85.5% [81.7–88.6]
Shaping 0.1, update 20	86.5% [82.8–89.5]
Shaping 0.1, update 30	86.8% [83.1–89.7]
Shaping 0.1, update 40	85.0% [81.2–88.2]

Every later evaluation therefore faced a different defender. Table 6 isolates the effect: with the same policy and the same harness, the Phase-A attacker falls from 86.7% to 14.0% when only the defender’s scoring changes. The 97% of Section 6.1 and the 8–33% of Sections 6.4 and 6.6 refer to two different opponents.

Sample size compounds the problem. With 12 games, the 95% interval on a win rate near 25% spans roughly ± 25 points. The heuristic baseline reported as 8% on 12 seeds is 15.3% on 300. Most differences in Tables 2 and 5 lie within one or two games, so they may be read as hypotheses rather than established effects. This applies, in particular, to the saturation of search and to the hierarchy law.

7.2 Where the attacker fails

Instrumenting 300 heuristic games shows that the attacker rarely reaches the stage in which objectives matter. It installs the bridge in only 40% of games, mostly in rounds 5–6, and the bridge survives to the end in 10%. In 68% of games no Israeli unit remains on the west bank, and 83% end with no objective held. The attrition clause of the victory rule (Israeli losses exceeding Israeli units on the west bank) accounts for 88% of Egyptian wins; with an empty west bank, however, that clause reduces to a failure to cross. Consistent with this reading, changing the casualty rule (spreading step losses instead of eliminating wounded units) or the post-combat advance (preferring objective cells) has no measurable effect: the attacker wins 17.0% under every variant ($n = 400$).

The same diagnosis could explain part of the search plateau. In a movement decision the attacker has about 126 legal moves. The policy places 74% of its mass on its top six, but the move the heuristic would make lies in that set in only 19% of decisions (24% in the top twelve). Because PUCT expands only the policy’s top six, it refines a narrow set of candidates, and further simulations cannot reach moves outside it.

7.3 Retraining against the current defender

If the attacker’s weakness is a distribution shift, retraining against the current defender should remove it. We warm-started PPO from the Phase-A MLP and trained against the expectimax heuristic, with and without potential-based shaping [3]. Table 7 reports validation in the full game loop on held-out seeds. Within 30 updates the attacker moves from 14% to about 87% against the very defender under which search had appeared to saturate near 33%. There is no sign of a structural ceiling.

Two caveats qualify this result. Training is unstable: the shaped run collapsed to 0% after update 50 and the unshaped run oscillated between 44% and 85%, so checkpoints are chosen by validation on

Table 8: Scoring mode by side (heuristic vs. heuristic, Israeli win rate, $n = 600$ per cell, Wilson 95% intervals).

	Egypt expectimax	Egypt fixed weights
Israel expectimax	11.3% [9.0–14.1]	23.8% [20.6–27.4]
Israel fixed weights	18.3% [15.4–21.6]	31.7% [28.1–35.5]

400 games. Most wins are also marginal (one objective held), with a mean graded score near 0.2, and the bridge is standing at the end of only about 1% of games.

7.4 Tuning the heuristics

The size of the expectimax effect suggests that the heuristic’s hand-set weights matter more than their casual choice implied. Table 8 crosses the two scoring modes by side. The effects are nearly additive: expectimax costs the attacker about seven points and gains the defender twelve to thirteen. The appropriate configuration is therefore asymmetric, fixed weights for the attacker and expectimax for the defender.

We then optimized each side’s weights with the cross-entropy method [8]: a diagonal Gaussian over normalized weights, 16 candidates per generation evaluated on 120 common seeds, fresh seeds at every generation, and a final validation on 1000 seeds not used in the optimization. For the attacker, against the default defender, the optimized weights raise the win rate from 18.6% [16.3–21.1] to 50.8% [47.7–53.9]. The largest change is a weaker pull toward distant targets (the distance weight falls from 10 to 4.5), which leaves relatively more weight on occupying targets, terrain and exposure, together with a slight bias of the bridge site toward the objectives. Explicit bridge doctrines did not help: we swept seven site rules and zero to three escorts that stay with the engineer after the bridge is built, and none improved on the default within the intervals. Guarding the engineer lowered the rate from 14% to 9%, and the installation rate stayed at 33–43% under every rule, whereas the retrained network installs the bridge in nearly every game. The bottleneck of the heuristic appears to lie in escorting the engineer to the canal rather than in choosing where to cross.

7.5 Search candidates and the defender’s prior

Adding the heuristic’s move to the candidate set tests the diagnosis of Section 7.2 directly. With the self-play network of Section 6.4 and 48 simulations, PUCT wins 23.5% [18.2–29.8] against the heuristic when it expands the policy’s top six, and 55.0% [48.1–61.7] when the heuristic’s move is added to them ($n = 200$ each). The plateau of Section 6.4 came from the candidate set, not from the value. With the retrained network the same change makes the attacker far more robust: its greedy policy wins none of 300 games against the retrained GNN or the tuned heuristic defender, whereas search with the union of candidates wins 16.7% and 48.3% of them ($n = 60$).

The defender shows the opposite effect. With the value of the retrained GNN, a PUCT defender whose candidates come from the heuristic lets the retrained attacker win 41.7% [30.1–54.3]; with the policy’s candidates, 3.3% [0.9–11.4]. The attacker has likely learned to exploit the doctrine, and a heuristic prior appears to re-import that doctrine into the search. A doctrinal move may thus help the side whose policy is weak and hurt the side whose policy is already strong.

7.6 Mutual exploitability and the population Nash

Each agent above was trained or tuned against a single opponent, and Table 9 illustrates how fragile this makes it. The retrained attacker wins 89.7% against the heuristic, 3.3% against the Phase-A GNN and none of 300 games against a GNN retrained against it. The tuned heuristic defender, which beats that attacker in all 1000 validation games, loses 63.3% of its games to the tuned heuristic attacker. The Phase-A MLP, weak against the current heuristic, beats the Phase-A GNN in 58% of games. No agent dominates.

We therefore summarize the population by its Nash equilibrium, computed by multiplicative weights [9] on the matrix of win rates. After one iteration of policy-space response oracles (PSRO [5]), in which each side adds a best response to the other side’s equilibrium mixture, the

Table 9: Israeli win rate (%) in the final population (rows: attackers; columns: defenders). $n = 300$ in cells with heuristic or network agents, whose 95% intervals lie within about ± 5 points, and 40–60 in cells involving search (about ± 13). Bold: support of the Nash equilibrium, in which Israel wins 17.0%.

Attacker	Heuristic		GNN			PUCT
	default	CEM	Phase A	retrained	PSRO	heur. prior
Heuristic, default	13.7	16.0	13.3	9.3	8.7	15.0
Heuristic, fixed weights	14.3	21.0	13.0	8.3	14.3	5.0
Heuristic, CEM vs. default	47.3	63.3	14.3	12.7	0.3	5.0
Heuristic, CEM vs. GNN	42.0	31.3	14.3	16.0	3.0	11.7
MLP, Phase A	21.0	0.3	58.0	4.0	10.0	21.7
MLP, retrained	89.7	0.0	3.3	0.0	9.3	41.7
MLP, PSRO	0.0	5.3	0.0	0.0	1.7	0.0
PUCT, retrained, union	95.0	48.3	31.7	16.7	25.0	77.5

attacker wins 17.0% at equilibrium (gap 0.5 points). Israel plays PUCT with the union of candidates in 98% of games, and Egypt plays the retrained GNN in 97%. The iteration barely moved this value, which was 16.9% before it. The attacker’s best response by PPO collapsed and lost even to the default heuristic. Its best response by CEM rose from 11.6% to 19.2% against the retrained GNN, but fell to 3% against the defender’s new best response, which in turn neutralized the tuned heuristic attacker (0.3%) at the cost of conceding more to search. The equilibrium agents are also the most robust in the table: against the retrained GNN no attacker exceeds 17%, and PUCT with the union of candidates is the strongest row in four of the six columns.

8 Discussion and Limitations

The game’s asymmetry is at once its richness and its challenge. Some lessons of the first part of the study are relative comparisons within a single protocol, and the audit does not bear on them directly: the best architecture depends on the role, naive self-play collapses the disadvantaged side, and the failure of small networks as attackers is one of capacity rather than inductive bias (Section 6.5). Others do not survive. The saturation of search and the hierarchy law rest on 12-game comparisons against an opponent that none of the agents had faced in training, and we keep them only as hypotheses. The ceiling itself was a distribution shift.

The audit adds three lessons. First, **the opponent is part of the result**. A change in the heuristic’s combat scoring moved a fixed policy from 87% to 14%, and a reweighting found by CEM moved a strong attacker to 0%; a single number against a single heuristic may say as much about the heuristic as about the agent. Second, **search is limited by what it is allowed to consider**. One doctrinal move among the candidates doubles the attacker’s search, whereas a heuristic prior for the defender re-imports the doctrine that its opponent has learned to exploit. Third, **balance is a property of a population**. Every agent trained against a single opponent was exploitable, and the equilibrium of the population, with the attacker near 17%, is likely the more informative summary of an asymmetric game.

Limitations: (i) the population is small (8 by 6) and the PSRO process stopped after one iteration, in which the attacker’s best response by PPO collapsed, so the equilibrium value could move with further iterations; (ii) cells involving search use 40–60 games, with intervals near ± 13 points; (iii) the PUCT transition model fixes the opponent to the heuristic and resolves chance by a single sample per edge; (iv) PPO fine-tuning is unstable, and checkpoints are chosen by validation on held-out games; (v) the AlphaZero loop is demonstrated, not scaled.

9 Conclusion

We documented an end-to-end RL system over an asymmetric historical wargame, from rules engine to in-browser inference, and then audited its evaluation. The audit reversed the study’s main negative claim. The attacker was not capped near a third of the games; we were evaluating it against a defender it had not faced, on samples too small to resolve the differences reported. Retraining,

tuning the heuristics and widening the candidates of search each lift the attacker substantially against a fixed opponent. None of these gains is robust, however: every agent trained against one opponent loses to some counter-strategy. The appropriate object of study is therefore the population and its equilibrium, in which the defender keeps the advantage and the attacker wins about 17% of games. The released artifact includes the evaluation arena, the optimizer and the population matrices needed to reproduce these results.

Reproducibility. Code, environment, trainer, and models live in `web/` (game) and `rl/` (RL). The PUCT search is `web/js/ai-mcts.js` (with the value head exposed in `ai-neural.js`); the AlphaZero loop lives in `rl/az_*` (self-play, training, evaluation, and the `az.py` orchestrator). The macro/LLM/ hierarchy controllers are `web/js/ai-macro.js` and `ai-hier.js` with `rl/llm-planner.mjs` (provider-agnostic, grounded) and `rl/skillopt.mjs` (text-space skill optimization with our simulator as verifier). The audit of Section 7 uses `rl/arena.mjs` (parallel evaluation with Wilson intervals), `rl/cem.mjs` (weight optimization), `rl/nash.mjs` and `rl/exp/` (population matrices, configurations and results); training against a fixed or mixed opponent is `rl/train.py` with `-oponente`. Commands in `rl/RL.md`; visual explanation in `rl/explicacao.html`; engine, rule-conformance, and inference-parity tests (logits and value) under `tests/` and `rl/test-parity.mjs`. Fixed seeds ensure reproducibility.

References

- [1] R. S. Sutton, A. G. Barto. *Reinforcement Learning: An Introduction*. 2nd ed., MIT Press.
- [2] J. Schulman et al. *Proximal Policy Optimization Algorithms*. 2017.
- [3] A. Y. Ng, D. Harada, S. Russell. *Policy Invariance under Reward Transformations: Theory and Application to Reward Shaping*. ICML, 1999.
- [4] D. Silver et al. *A general reinforcement learning algorithm that masters chess, shogi, and Go through self-play*. Science, 2018.
- [5] M. Lanctot et al. *A Unified Game-Theoretic Approach to Multiagent Reinforcement Learning*. NeurIPS, 2017.
- [6] J. Schulman et al. *High-Dimensional Continuous Control Using Generalized Advantage Estimation*. ICLR, 2016.
- [7] E. B. Wilson. *Probable Inference, the Law of Succession, and Statistical Inference*. Journal of the American Statistical Association 22:209–212, 1927.
- [8] P.-T. de Boer, D. P. Kroese, S. Mannor, R. Y. Rubinstein. *A Tutorial on the Cross-Entropy Method*. Annals of Operations Research 134:19–67, 2005.
- [9] Y. Freund, R. E. Schapire. *Adaptive Game Playing Using Multiplicative Weights*. Games and Economic Behavior 29:79–103, 1999.